


Executive Summary

- Framing: task-conditioned affordance candidate ranking
- Partial semantic grounding — clean semantic success (unmasked): 8/9
- Mean entropy (unmasked/masked): 1.057 / 0.793



Problem & Architecture


- Affordance ranking under noisy supervision
- Lightweight Ranker + task-conditioning + compatibility masking
- Metrics: top1/top3, entropy, compatible rank, reject behavior



Calibration-vs-Ranking tradeoff (masking ablation)

Entropy and Incompatible-score Distributions





Reject behavior matrix by case type

Exemplars: best (top) / worst (bottom)

4 / no_valid_affordance
baseball bat (1.000)



4 / no_valid_affordance
snowboard (0.999)



11 / distractor_heavy
dining table (0.375)



11 / ambiguous_affordance
fork (0.142)





Masking false-accept exemplar — demonstrates shortcut priors, calibration distortion, ambiguity sensitivity, and semantic contamination
Image: COCO_train2014_000000131498.jpg | task_id: 1 | case_type: ambiguous_affordance